



Slides by Andrew DeOrio,
Jeff Ringenberg
and Brian Noble

EECS 280

Programming and Introductory Data Structures

Container ADTs

Review: Data Abstraction

- **Data abstraction** lets us separate *what* a type is (and what it can do) from *how* the type is implemented
- In C++, we use a `class` to implement data abstraction
 - We can create an Abstract Data Type (ADT) using a `class`
- ADTs let us model complex phenomena
 - More complex than built-in data types like `int`, `double`, etc.
- ADTs make programs easier to maintain and modify
 - You can change the implementation and no users of the type can tell

Review: Information Hiding and Encapsulation

Information Hiding

- Protect and hide our code from other code that that uses it

Encapsulation

- Keeping data and relevant functions together

Containers

- The purpose of a *container* is to hold other objects

- For example, an array can hold integers:

```
int array[10];  
a[4] = 517;
```

- What if we want more features in our container? Say, a set of integers.
- This is a “set” in the mathematical sense: A collection of zero or more integers, with no duplicates.

Set abstraction

- Let's create an abstraction that holds a set of integers
- There are four operations on this set that we will define:
 - Insert a value into the set
 - Remove a value from the set
 - Query to see if a value is in the set
 - Count the number of elements in the set

IntSet Declaration

```
class IntSet {  
    public:  
        IntSet();  
        void insert(int v);  
        void remove(int v);  
        bool query(int v) const;  
        int  size() const;  
  
        //...  
};
```

IntSet Declaration

```
class IntSet {  
    //...  
  
    //EFFECTS: returns true if v is in set,  
    //false otherwise  
    bool query(int v) const;  
  
    //EFFECTS: returns |set|  
    int size() const;  
  
    //...  
};
```

IntSet Representation

```
class IntSet {  
    public:  
        IntSet();  
        void insert(int v);  
        void remove(int v);  
        bool query(int v) const;  
        int size() const;  
};
```

- The class is incomplete because we haven't chosen a representation for sets
- Choosing a representation involves two things:
 1. Deciding what concrete data elements will be used to represent the values of the set
 2. Providing an implementation for each method

IntSet Representation

```
class IntSet {  
    public:  
        IntSet();  
        void insert(int v);  
        void remove(int v);  
        bool query(int v) const;  
        int size() const;  
};
```

- Despite not having a representation for a set, the (incomplete) declaration is all that a **customer** of the IntSet abstraction needs to know since it has:
 - The general overview of the ADT.
 - The specification of each method.

IntSet Representation

- Start with a representation invariant for the set:
- Use an array
- Represent a set of size N as an unordered array of integers with no duplicates, stored in the first N slots of the array.
- `int elts_size`: equal to the number of elements currently in the array.

IntSet Representation

- Since this is an array, and arrays have maximum sizes, we have to choose a maximum size

```
int ELTS_CAPACITY = 100;
```

- 100 is arbitrary, but soon we'll get to dynamic memory ☺

- Include the size in the OVERVIEW

```
//OVERVIEW: mutable set of ints,  
|set|<=ELTS_CAPACITY
```

- Account for full sets in the REQUIRES clause for insert:

```
//REQUIRES: set is not full  
//MODIFIES: this  
//EFFECTS: set=set+{v}  
void insert(int v);
```

IntSet::size()

```
class IntSet {
    //OVERVIEW: mutable set of ints, |set| <=ELTS_CAPACITY
public:
    void insert(int v);
    void remove(int v);
    bool query(int v) const;
    int size() const;
    static const int ELTS_CAPACITY = 100;
private:
    int      elts[ELTS_CAPACITY];
    int      elts_size;
};
```

static Member Variables

```
class IntSet {  
public:  
    static const int ELTS_CAPACITY = 100;  
};
```

- `static` means "there's only one"
 - All instances of the class share this member variable
 - A lot like a global variable
- `const` means "you can't change this"

static is Confusing!

```
static int add(int a, int b) {  
    return a + b;  
}
```

- Static function: visible only inside one file (internal linkage)

```
class IntSet {  
public:  
    static const int ELTS_CAPACITY = 100;  
};
```

- Static member variable: all instances of the `IntSet` class share this member variable

IntSet::size()

```
class IntSet {  
    public:  
        void insert(int v);  
        void remove(int v);  
        bool query(int v) const;  
        int size() const;  
        static const int ELTS_CAPACITY = 100;  
    private:  
        int        elts[ELTS_CAPACITY];  
        int        elts_size;  
};
```

```
int IntSet::size() const {  
    return elts_size;  
}
```

Given this representation, and the representation invariants, we can write the methods.

Because our rep invariant says that `elts_size` is always the size of the set, we can return it directly

Searching IntSet

- Next, consider the three final member functions:
 - `query`: *search* the array looking for a specific number
 - `remove`: *search* the array for a number; if it exists, remove it
 - `insert`: *search* the array for a number; if it doesn't exist, add it
- All three of these have *search* in common.
- One might be tempted to just write `insert` and `remove` in terms of `query`, but `remove` won't work that way.
- `Query` only tells us **whether** the element exists, not **where** – we need one more method

IntSet::indexOf()

```
class IntSet {
public:
    void insert(int v);
    void remove(int v);
    bool query(int v) const;
    int size() const;
    static const int ELTS_CAPACITY = 100;
private:
    int          elts[ELTS_CAPACITY];
    int          elts_size;
    int indexOf(int v) const;
    // EFFECTS: returns the index of
    //          v if it exists in the
    //          array, ELTS_CAPACITY otherwise
};
```

This member function must be **private**. This is because it exposes details about the concrete representation. It is inappropriate to expose these details to a client of this class.

IntSet::indexOf()

```
class IntSet {
public:
    //...
    static const int ELTS_CAPACITY = 100;
private:
    int          elts[ELTS_CAPACITY];
    int          elts_size;
    int indexOf(int v) const;
};

int IntSet::indexOf(int v) const {
    for (int i = 0; i < elts_size; ++i) {
        if (elts[i] == v) return i;
    }
    return ELTS_CAPACITY;
}
```

IntSet::query()

```
class IntSet { // OVERVIEW omitted for space
public:
    bool query(int v) const;
    //...
    static const int ELTS_CAPACITY = 100;
private:
    int          elts[ELTS_CAPACITY];
    int          elts_size;
    int indexOf(int v) const;
};
```

With indexOf, query is easy!

```
bool IntSet::query(int v) const {
    return indexOf(v) != ELTS_CAPACITY;
}
```

IntSet::insert()

- The code for `insert` is not much more difficult than `query`:
 1. First look for the `indexOf` the element to insert
 2. If it doesn't exist, we need to add this element to the “end” of the array
 3. Using `elts_size`, the current “end” is `elts[elts_size-1]`
 4. Place the element in the next slot and update `elts_size`

IntSet::insert()

```
class IntSet {
public:
    void insert(int v);
    bool query(int v) const;
    //...
    static const int ELTS_CAPACITY = 100;
private:
    int      elts[ELTS_CAPACITY];
    int      elts_size;
    int indexOf(int v) const;
};

void IntSet::insert(int v) {
    if (query(v)) return;
    assert(elts_size < ELTS_CAPACITY); //REQUIRES!
    elts[elts_size++] = v;
}
```

Take a couple minutes
to figure this out.

IntSet::remove ()

- `remove` works similarly to `insert`.
 1. If the element (called the `victim`) is in the array we have to remove it leaving a "hole" in the array
 2. Instead of moving each element after the `victim` to the left by one position, pick up the current "last" element and move it to the hole
- This again breaks the invariant on `elts_size`, so we must fix it

IntSet::remove ()

```
void IntSet::remove(int v) {  
    int victim = indexOf(v);  
    if (victim == ELTS_CAPACITY) return; //not found  
    elts[victim] = elts[--elts_size];  
}
```

Take a couple minutes
to figure this out

Initialization Exercise

- Consider the newly-created set

```
int main() {  
    IntSet s;  
}
```

- **Problem:** On creation, `s`'s data members are uninitialized!
- This means that the value of `elts_size` could be anything, but our representational invariant says it must be zero!
- Exercise: implement the `IntSet` default constructor

Initialization Exercise

```
IntSet::IntSet()  
    : elts_size(0) {}
```

Exercise

- Add a `print()` member function
- "Promise" not to modify any member variables

```
class IntSet {
public:
    void insert(int v);
    void remove(int v);
    bool query(int v) const;
    int  size() const;
    static const int ELTS_CAPACITY = 100;
private:
    int      elts[ELTS_CAPACITY];
    int      elts_size;
};
```

Exercise

```
class IntSet {  
    //...  
    //EFFECTS: prints set  
    void print() const;  
};  
  
void IntSet::print() const {  
    cout << "{ ";  
    for (int i=0; i<elts_size; ++i)  
        cout << elts[i] << " ";  
    cout << "}" << endl;  
}
```

Using IntSet

```
int main () {  
    IntSet is;  
    is.insert(7);  
    is.insert(4);  
    is.insert(7);  
    is.print();  
    is.remove(7);  
    is.print();  
}
```

```
./a.out  
{ 7 4 }  
{ 4 }
```

IntSet efficiency

```
int IntSet::indexOf(int v) const {  
    for (int i = 0; i < elts_size; ++i) {  
        if (elts[i] == v) return i;  
    }  
    return ELTS_CAPACITY;  
}
```

- **Question:** How many elements of the array must `indexOf` examine in the worst case if there are 10 elements? If there are 90 elements?

IntSet efficiency

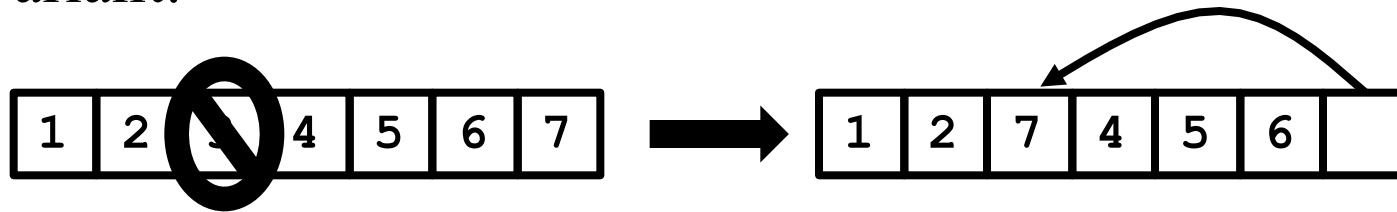
- We say the time for `indexOf` grows *linearly* with the size of the set.
- If there are N elements in the set, we have to examine all N of them in the worst case. For large sets that perform lots of queries, this might be too expensive.
- Luckily, we can replace this implementation with a different one that can be more efficient. The only change we need to make is to the representation – the abstraction can stay precisely the same.

Improving IntSet efficiency

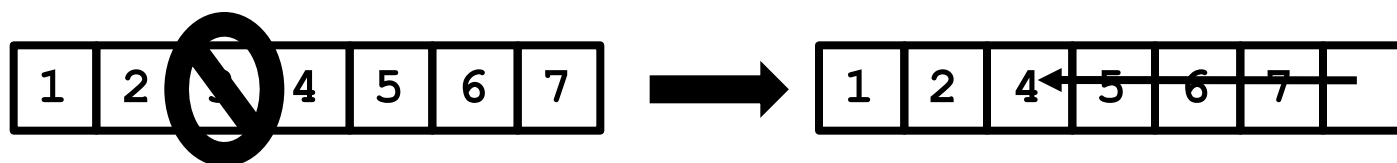
- Still use an array (of 100 elements) to store the elements of the set and the values will still occupy the first `elts_size` slots.
- *However, now we'll keep the elements in sorted order*
- The constructor and size methods don't need to change at all since they just use the `elts_size` field.
- `query` also doesn't need to change. If the index exists in the array's legal bounds, then it's there.

Improving IntSet efficiency

- However, the others all do need to change. We'll start with the easiest one: `remove()`
- Recall the old version that moved the last element from the end to somewhere in the middle, this will break the new “sorted” invariant.



- Instead of doing a swap, we have to "squish" the array together to cover up the hole.



Improving IntSet efficiency

- How are we going to do the “squish”?
- Move the element next to the hole to the left leaving a new hole
- Keep moving elements until the hole is “off the end” of the elements

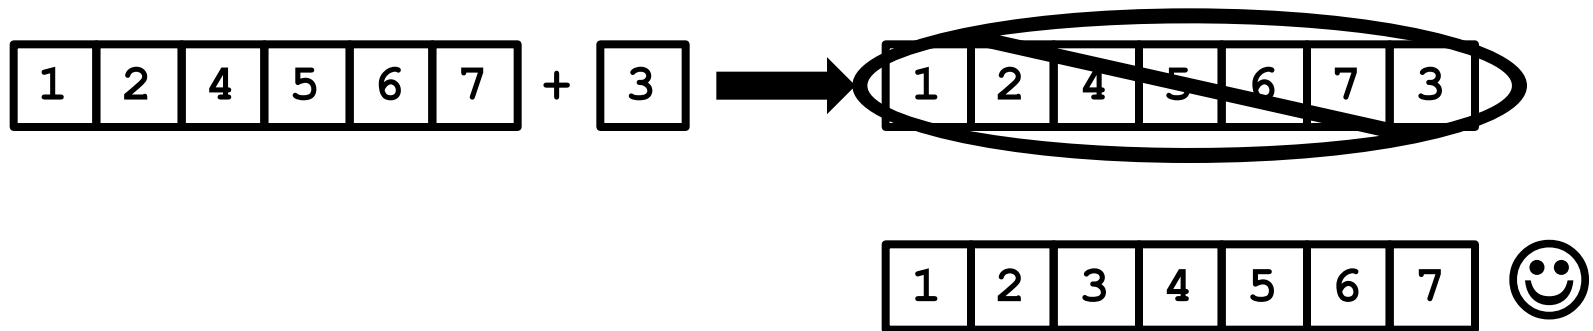
Improving IntSet efficiency

```
void IntSet::remove(int v) {  
    int gap = indexOf(v);  
    if (gap == ELTS_CAPACITY) return; //not found  
    --elts_size; //one less element  
  
    while (gap < elts_size) {  
        //there are elts to our right  
        elts[gap] = elts[gap+1];  
        ++gap;  
    }  
}
```

Take a couple minutes
to figure this out

Improving IntSet efficiency

- We also have to change `insert` since it currently just places the new element at the end of the array. This also will break the new “sorted” invariant.



Improving IntSet efficiency

- How are we going to do the insert?
- Start by moving the last element to the right by one position.
- Repeat this process until the correct location is found to insert the new element.
- Stop if the start of the array is reached or the element is sorted.
- We'll need a new loop variable to track this movement called `cand(idate)`.
- It's invariant is that it always points to the next element that might have to move to the right.

Improving IntSet efficiency

```
void IntSet::insert(int v) {  
    if (indexOf(v) != ELTS_CAPACITY) return; //present!  
    assert(elts_size < ELTS_CAPACITY); //REQUIRES  
    int cand = elts_size-1; // largest element  
    while ((cand >= 0) && elts[cand] > v) {  
        elts[cand+1] = elts[cand];  
        --cand;  
    }  
    // Now, cand points to the left of the "gap"  
    elts[cand+1] = v;  
    ++elts_size; // repair invariant  
}
```

Take a couple minutes
to figure this out

Improving IntSet efficiency

```
void IntSet::insert(int v) {  
    if (indexOf(v) != ELTS_CAPACITY) return; //present!  
    assert(elts_size <= ELTS_CAPACITY) //enough room?  
    int cand = elts_size-1; // largest element  
    while ((cand >= 0) && elts[cand] > v) {  
        elts[cand+1] = elts[cand];  
        --cand;  
    }  
    // Now, cand points to the largest element less than v  
    elts[cand+1] = v;  
    ++elts_size; // repair invariant  
}
```

We are using the "short-circuit" property of &&. If cand is not greater than or equal to zero, we never evaluate the right-hand clause.

Improving IntSet efficiency

- **Question:** Do we have to change `indexOf`?

```
int IntSet::indexOf(int v) const {  
    for (int i = 0; i < elts_size; ++i) {  
        if (elts[i] == v) return i;  
    }  
    return ELTS_CAPACITY;  
}
```

Improving IntSet efficiency

- **Question:** Do we have to change `indexOf`?
- **Answer:** No, but it can be made more efficient with the new representation

```
int IntSet::indexOf(int v) const {  
    for (int i = 0; i < elts_size; ++i) {  
        if (elts[i] == v) return i;  
    }  
    return ELTS_CAPACITY;  
}
```


Improving IntSet efficiency

- Suppose we are looking for x .
- Compare x against the **middle** element of the array and there are three possibilities:
 1. x is equal to the middle element.
 2. x is less than the element.
 3. x is greater than the element.
- If it's case 1, we're done.
- If it's case 2, then if x is in the array, it must be to the **left** of the middle element
- If it's case 3, then if x is in the array, it must be to the **right** of the middle element.

Improving IntSet efficiency

- Compare `foo` against the **middle** element of the array and there are three possibilities:
 1. `foo` is equal to the middle element.
 2. `foo` is less than the element.
 3. `foo` is greater than the element.
- The comparison with the middle element eliminates at least half of the array from consideration! Then, we repeat the same thing over again.
- You could write this "repetition" as either a tail-recursive program or an iterative one. Most programmers find the iterative version more natural, so we'll write it iteratively, too.

Improving IntSet efficiency

- First, we need to find the “bounds” of the array.
- The “leftmost” element is always zero, but the “rightmost” element is `elts_size-1`.

```
int IntSet::indexOf(int v) const {  
    int left = 0;  
    int right = elts_size-1;  
    ...  
}
```

Improving IntSet efficiency

- It's possible that the segment we are examining is empty and we return `ELTS_CAPACITY` since the element is missing.
- A nonempty array has at least one element in it, so `right` is at least as large as `left` (`right >= left`).

```
int IntSet::indexOf(int v) const {  
    int left = 0;  
    int right = elts_size-1;  
    while (right >= left) {  
        ...  
    }  
    return ELTS_CAPACITY;  
}
```

Improving IntSet efficiency

- Next, find the "middle" element. We do this by finding out the size of our segment ($\text{right} - \text{left} + 1$), then divide it by two, and add it to left.

```
int IntSet::indexOf(int v) const {
    int left = 0;
    int right = elts_size-1;
    while (right >= left) {
        int size = right - left + 1;
        int middle = left + size/2;
        ...
    }
    return ELTS_CAPACITY;
}
```

Improving IntSet efficiency

- Next, find the "middle" element. We do this by finding out the size of our segment ($\text{right} - \text{left} + 1$), then divide it by two, and add it to left.

```
int IntSet::indexOf(int v) const {  
    int left = 0;  
    int right = elts_size-1;  
    while (right >= left) {  
        int size = right - left + 1;  
        int middle = left + size/2;  
        ...  
    }  
    return ELTS_CAPACITY;  
}
```

If there is an odd number of elements, this will be the "true" middle. If there are an even number, it will be the element to the "right" of true middle.

Improving IntSet efficiency

- Then, we compare against the middle element.
- If that's the one we are looking for, we are done.

```
int IntSet::indexOf(int v) const {  
    int left = 0;  
    int right = elts_size-1;  
    while (right >= left) {  
        int size = right - left + 1;  
        int middle = left + size/2;  
        if (elts[middle] == v) return middle;  
        ...  
    }  
    return ELTS_CAPACITY;  
}
```

Improving IntSet efficiency

- If we are not looking at the element, the true element (if it exists) must be in either the “smaller” half or the “larger” half.
- If it is in the smaller half, then we can eliminate all elements at index `middle` and higher, so we move “right” to `middle-1`.
- Likewise, if it would be in the larger half, we move “left” to `middle+1`, and we continue looking.

Improving IntSet efficiency

```
int IntSet::indexOf(int v) const {
    int left = 0;
    int right = elts_size-1;
    while (right >= left) {
        int size = right - left + 1;
        int middle = left + size/2;
        if (elts[middle] == v)
            return middle;
        else if (elts[middle] < v)
            left = middle+1;
        else
            right = middle-1;
    }
    return ELTS_CAPACITY;
}
```

Take a couple minutes to figure this out. Try using the inputs: $v=3$, $v=8$, and $v=-1$ with the array below.

0	1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---	---

Improving IntSet efficiency

- Since you eliminate half of the array with each comparison, this is a much more efficient.
- If the array has N elements, you'll need “about” $\log_2(N)$ comparisons to search it.
- This is really cool, because $\log_2(100)$ is less than 7 – so, we need only 7 comparisons in the **worst case**.
- Also, if you double the size of the array, you need only one extra comparison to do the search.
- This is called a *binary search*

Improving IntSet efficiency

- `insert` and `remove` are still linear, because they may have to "swap" an element to the beginning/end of the array.
- Here is the summary of asymptotic performance of each function:

	<u>Unsorted</u>	<u>Sorted</u>
<code>insert</code>	$O(N)$	$O(N)$
<code>remove</code>	$O(N)$	$O(N)$
<code>query</code>	$O(N)$	$O(\log N)$

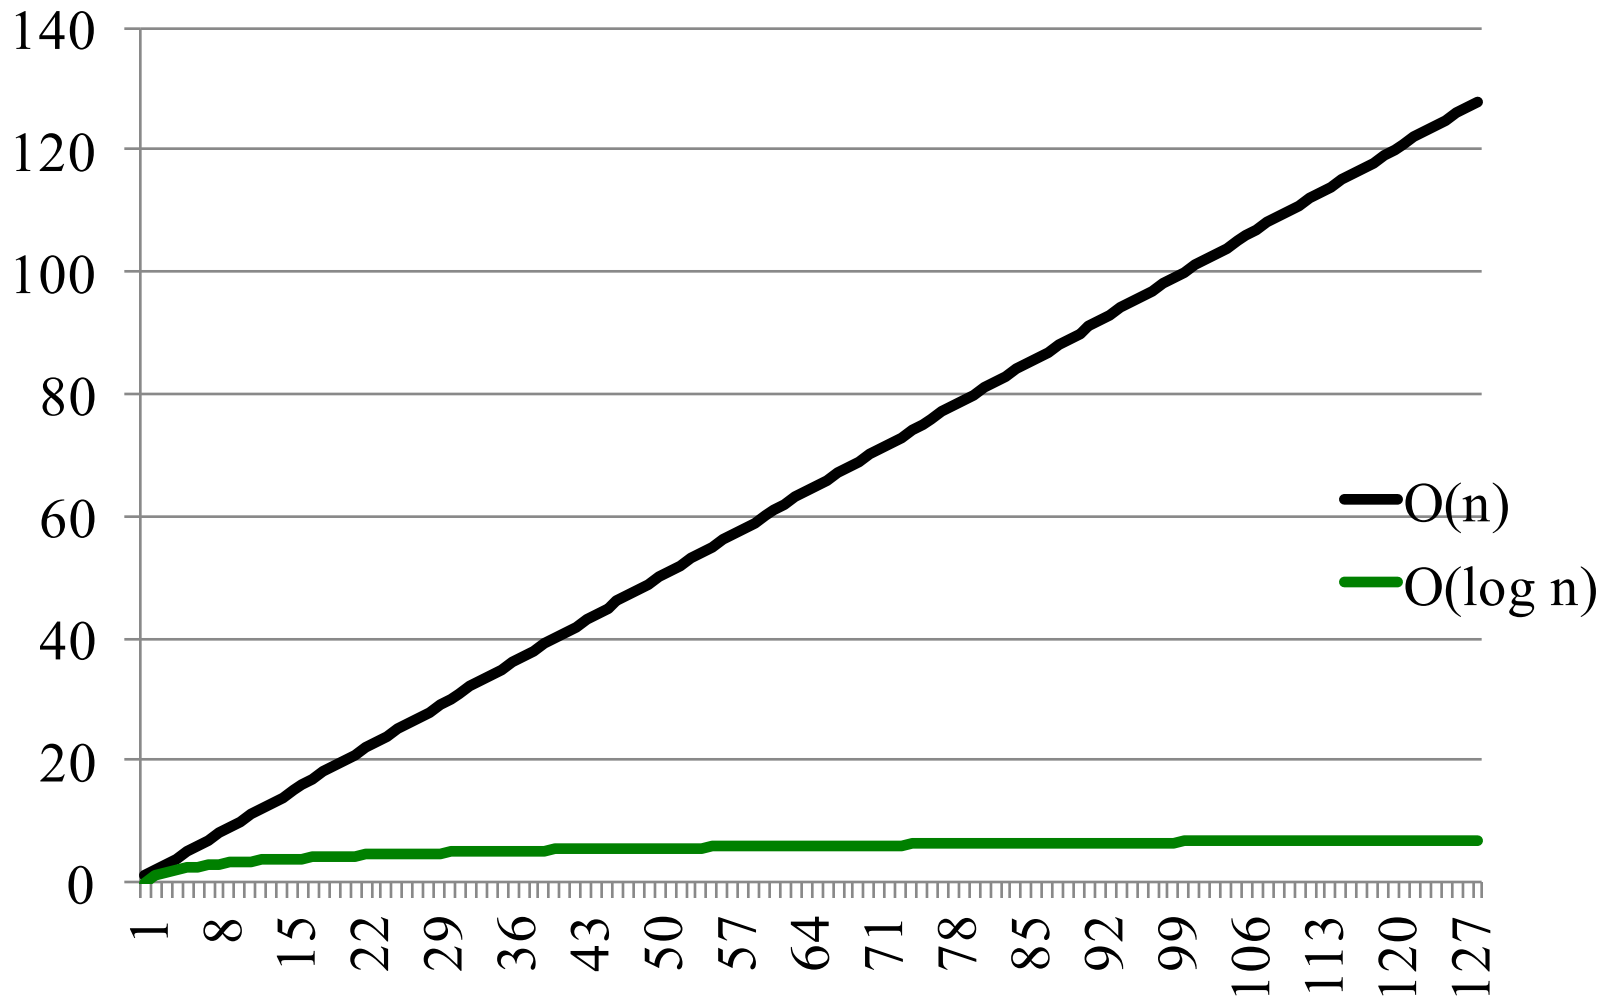
All the unsorted versions still require the unsorted `indexOf()` which is $O(N)$

Improving IntSet efficiency

	<u>Unsorted</u>	<u>Sorted</u>
<code>insert</code>	$O(N)$	$O(N)$
<code>remove</code>	$O(N)$	$O(N)$
<code>query</code>	$O(N)$	$O(\log N)$

- If you are going to do more searching than inserting/removing, you should use the "sorted array" version, because `query` is faster there.
- However, if `query` is relatively rare, you may as well use the "unsorted" version. It's "about the same as" the sorted version for `insert` and `remove`, but it's MUCH simpler!

Log(n) run time. So what?



Abstraction

We modified `IntSet`. Do we need to change our `main()` program? Why?

```
int main () {  
    IntSet is;  
    is.insert(7);  
    is.insert(4);  
    is.insert(7);  
    is.print();  
    is.remove(7);  
    is.print();  
}
```

Abstraction

We modified `InSet`. Do we need to change our `main()` program? Why?

No modification needed, because of data abstraction