



Slides by Andrew DeOrio,
Jeff Ringenberg and
Brian Noble

EECS 280

Programming and Introductory Data Structures

Memory Models

Static, fixed-size structures

- So far, the data structures we've built have all had room for "at most N " elements:
- The various `IntSet` implementations could have at most 100 distinct integers
- We could extend these sizes to larger ones, but we really only know how to create *static, fixed-sized* structures.

Static, fixed-size structures

- Sometimes, the process you are modeling has a physical limit, which makes a static, fixed-sized structure a reasonable choice.
- For example, a deck of cards has 52 individual cards in it, so this is a reasonable limitation.
- However, there is no meaningful sense in which a "set of integers" is limited to some particular number of elements.
- No matter how big you make the set's capacity, an application that needs more will eventually come along.
- Consider the `list_t` type from the second project:
 - The type imposed no limits on how large a list could grow

Global and local variables

- For the variables we have used so far you need to know two things at compile time (statically)
 - The size (or number)
 - The lifetime (when it will be created and destroyed)
- There have been two classes of such variables:
 1. Global variables
 2. Local variables

Global variables

- Global variables are defined anywhere outside of a function definition
- Space is set aside for these variables before the program begins execution, and is reserved for them until the program completes
- This space is reserved at compile time

Global variable examples

```
const int SIZE=10;  
int main() {  
    //...  
}
```

```
static Triangle  
    g_triangle;  
Shape * ask_user() {  
    //...  
    return &g_triangle;  
}
```

```
int sum_cur=0; //bad idea  
int sum(list_t list) {  
    // ...  
}
```

```
int prod_cur=0; //bad idea  
int product(list_t list) {  
    // ...  
}
```

Local variables

- Local variables are any variable defined within a block
- This includes function arguments, which act as if they were defined in the outermost block of the function
- Space is set aside for these variables when the relevant block is entered, and is reserved for them until the block is exited
- This space is reserved at run time, but the size is known to the compiler
- Since the compiler must know how big all of these variables will be, it is static information, and must be declared by the programmer

Local variable examples

```
int sum(list_t list) {  
    //...  
}
```

```
int sum(int *array, int size) {  
    int sum=0;  
    //...  
}
```

```
int main(int argc, char* argv[]) {  
    IntSet i;  
    for (int i=0; i<10; ++i) {  
        //...  
    }  
}
```


Dynamic variables

- There is a third type of variable, called a *dynamic variable*
- It is dynamic because
 - Size (or number) is determined at runtime
 - When it will be created and destroyed is determined at runtime
- In other words, you (the programmer) get to decide how big a dynamic variable is, when it is created, and when it is destroyed

new

- Create dynamic variables using new

```
int main() {  
    int *p = new int;  
}
```

- This creates new space for an integer, and returns a pointer to that space, assigning it to p
- The initial value is undefined
- Use initializer syntax to assign an initial value

```
int main() {  
    int *p = new int(5);  
}
```

Exercise


```
int main() {  
    int *p = new int(5);  
}
```

- How many variables are in this program?
- Mark each as global, local or dynamic

Exercise

```
int main() {  
    int *p = new int(5);  
}
```

- How many variables are in this program? **2**
 1. `int *p` is a local variable of type "pointer to integer"
 2. `int(5)` is a dynamic variable of type "integer"
- This “thing pointed to by `p`” is a dynamically-allocated piece of memory, and lives “somewhere else”.
- It does not have a name of its own, but is “the integer pointed to by `p`”.



delete

- Destroy dynamic variables using `delete`

```
int *p = new int;  
//do something with p  
delete p;
```

- Releases the claim on the space previously used by the `int`
- Space can be recycled later by `new`

delete pitfalls

- You can only delete a dynamic variable *once*

```
int *p = new int;  
//do something with p  
delete p;  
delete p; //Error!
```

```
*** glibc detected ***  
./a.out: double free or  
corruption (fasttop):  
0x00000000007ac010 ***  
...  
Aborted (core dumped)
```

delete pitfalls

- `delete 0` (AKA `NULL`) does nothing

```
int *p = new int;  
//do something with p  
delete p; p=0;  
delete p; //OK
```

delete pitfalls

- Ordinary objects can be destroyed by `delete`, *but only if they were created by `new`*

```
int i = 0; //local variable
```

```
int *p=&i; //pointer to local variable
```

```
delete p; //undefined! (likely a runtime error)
```


Size of dynamic variables

- With dynamic arrays, we can choose the size at runtime

```
//ask user to enter integer size  
int size = get_size_from_user();
```

```
int *p = new int[size];  
//do something with p ...  
delete[] p;
```

- Note the different syntax with `delete`
- We will talk about dynamic arrays in more depth next time

Lifetime of dynamic variables

- What's the problem?

```
//EFFECTS: allocates an array of specified size
// and initializes each element to zero
int * get_zero_array(int size) {
    int array[size];
    for (int i=0; i<size; ++i) array[i] = 0;
    return array;
}

int main() {
    int *a = get_zero_array(3);
    cout << a[0] << endl;
}
```

Lifetime of dynamic variables

- The lifetime of an object is completely under the control of the programmer – it lives until it is explicitly destroyed
- You can create a variable in one function and use it in another

```
//EFFECTS: allocates an array of specified size
//  and initializes each element to zero
int * get_zero_array(int size) {
    int *array = new int[size];
    for (int i=0; i<size; ++i) array[i] = 0;
    return array;
}
```

Lifetime of dynamic variables

```
//EFFECTS: allocates an array of specified size  
// and initializes each element to zero
```

```
int * get_zero_array(int size) {  
    int *array = new int[size];  
    for (int i=0; i<size; ++i) array[i] = 0;  
    return array;  
}
```

```
int main() {  
    int *a = get_zero_array(10);  
    //use a  
    delete[] a; a=0;  
}
```


Call `new` in one
function, and call
`delete` in another
function

Memory leaks

- Dynamic variables live until the programmer destroys them using `delete`
- This is true even if you "forget" the pointer to the object.


```
int main() {  
    int *p1 = new int(1);  
    int *p2 = new int(2);  
    p1 = p2;  
}
```

- This leaves us with:



Memory leaks

```
int main() {  
    int *p1 = new int(1);  
    int *p2 = new int(2);  
    p1 = p2;  
}
```




- Two pointers point to the object "2", and **none** to the object "1"
- There is no way to release the memory occupied by "1"

Memory leaks

- Note there is an important difference between the lifetime of a pointer variable and the lifetime of the object it points to!

```
int main() {  
    int *p1 = new int(1);  
    int *p2 = new int(2);  
    p1 = p2;  
}
```



- In this example, exiting the block that defines `p1` causes the local object `p1` to vanish, but the dynamic object it points to remains
- This leaves us with an allocated dynamic object that we have no means of reclaiming called a memory leak
- If memory leaks occur often enough, your program may reach a point where it can no longer allocate new dynamic objects

Why use dynamic variables?


1. If you want to manage the lifetime of the variable yourself
 - Not stack, Not global, but something in between
2. Many parts of the code refer to an object, but only want one copy in memory
 - Normally, if many parts refer to the same object, those parts will be in different scopes

The heap

- The space for objects created via `new` comes from a location in memory called the heap.
- To describe the heap, we first have to revisit the memory model used by a "typical" C++ process:
- A running program has an "address space", a collection of memory locations that are accessible to it
- An address space is private to a running program – no other running program can access/modify it

Program segments

- There are typically five parts, called segments, in an address space:




Address MAX

Address 0

Text segment

- The code comprising a compiled program goes in the text segment.
- It is at a very low address, but typically not at address zero.




Address MAX

Address 0

Globals and heap segments

- Immediately above the *text* section, the compiler allocates space for any global variables, and initializes them, when necessary.
- When dynamic variables are allocated with `new`, they come from the *heap*, which grows upward.

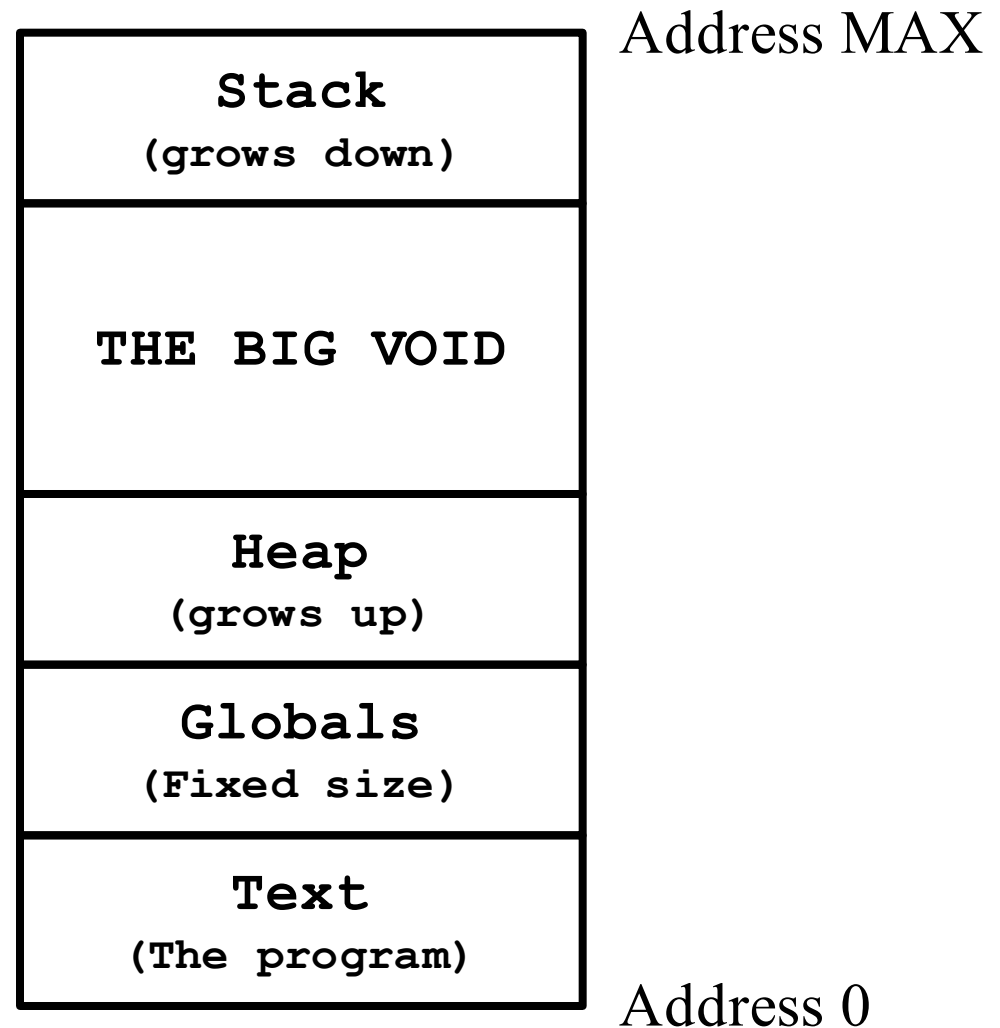


Address MAX

Address 0


Stack segment and the Big Void

- When functions are called, stack frames are created on the *stack*, which grows downward.
- Since we don't know how big either of these will get, we keep a big hole in between the two called THE BIG VOID.



The little void

- Most systems also reserve the first few thousand addresses starting at zero for another void.



Address MAX

Address 0


Global vs. local vs. dynamic

	Global	Local	Dynamic
Where in code?	Outside function	Inside function (block) or args	Anywhere you use <code>new</code>
When created	Beginning of program	Beginning of function (block)	You call <code>new</code>
When destroyed	End of program	End of function (block)	You call <code>delete</code>
Size	static	static	dynamic
Location	Globals	Stack	Heap

Stack and heap example

```
int main() {  
    int *p;  
}
```


- Local variable goes on the stack
- Undefined value
 - An address, since p is a pointer
- Could point to a random, part of memory



Stack and heap example

```
int main() {  
    int *p;  
    new int(100);  
}
```


- Dynamic memory goes on the heap
- Leaks `sizeof(int)` bytes of memory
- No way to find `int(100)` and delete it!



Stack and heap example

```
int main() {  
    int *p = new int(100);  
}
```

- Good habit: always initialize variables




Stack and heap example

```
int main() {  
    int *p = new int(100);  
    cout << *p << endl;  
}
```

100


- Access dynamic variable using pointer



Stack and heap example

```
int main() {  
    int *p = new int(100);  
    cout << *p << endl;  
}
```


- Problem: memory leak!



Stack and heap example

```
int main() {  
    int *p = new int(100);  
    delete p;  
    cout << *p << endl;  
}
```


- Problem: memory leak!
- Fixed
- Problem: reading dynamic variable after delete



Stack and heap example

```
int main() {  
    int *p = new int(100);  
    delete p; p=0;  
    cout << *p << endl;  
}
```


- Problem: reading dynamic variable after `delete`
- Fixed
- Problem: dereferencing a 0 (AKA NULL) pointer!



Stack and heap example

```
int main() {  
    int *p = new int(100);  
    delete p; p=0;  
    assert(p) ;  
    cout << *p << endl;  
}
```


- Problem: dereferencing a 0 (AKA NULL) pointer!
- `assert()` can help identify problem before it causes a SEGFAULT



Stack and heap example

```
int main() {  
    int *p = new int(100);  
    assert(p);  
    cout << *p << endl;  
    delete p; p=0;  
}
```

- Bugs fixed



NULL vs. 0

```
int *ptr = NULL;
```

- C style
- No library needed
- `int *ptr=0;`
 - Works, but `ptr=NULL` is preferred


```
int *ptr = 0;
```

- C++ style
- `int *ptr=NULL;`
 - *error: 'NULL' was not declared in this scope*
- `#include <cstddef>`
`int *ptr=NULL;`
 - Works, but `ptr=0` is preferred

Stack and heap exercise

```
int i = 42;  
int *p = &i;  
delete p; p=0;
```


- Draw the stack and the heap
- What is wrong with this code?



Stack and heap exercise

```
int i=4;  
int *p = new int(17);  
i = *p;  
delete p; p=0;
```


- Draw the stack and the heap
- How much memory is leaked?



Stack and heap exercise

```
int *p = new int(100);  
int *q = p;  
delete q; q=0;  
cout << *p << endl;
```


- Draw the stack and the heap
- What does this print?
- How much memory is leaked?



Stack and heap exercise

```
int *p = new int(100);  
int *q = new int(42);  
q=p;  
delete q; q=0;
```

- Draw the stack and the heap
- How much memory is leaked?



Classes and dynamic memory

- When you create instances of classes, their constructors are called, just as if it were created the "normal" way.

```
IntSet *isp = new IntSet;
```

1. Allocate enough space on the heap to hold an `IntSet`.
 - An array of 100 integers (`elts`)
 - One extra integer (`elts_size`) to hold its size
2. Call the constructor `IntSet::IntSet()` on this new object

Classes and dynamic memory

- We can also destroy instances of ADTs that were created by new:

```
IntSet *isp = new IntSet;  
delete isp; isp=0;
```

Exercise: allocating classes

```
IntSet *isp =  
    new IntSet;  
delete isp; isp=0;
```

- Draw the stack and heap
- Assume no virtual functions

